

Pen down the limitations of MapReduce

1. Processing speed

In Hadoop, with a parallel and distributed algorithm, MapReduce process large data sets. MapReduce algorithm contains two important tasks: Map and Reduce and, MapReduce require lot of time to perform these tasks thereby increasing latency. Data is distributed and processed over the cluster in MapReduce.

2. Data processing

Hadoop MapReduce is designed for Batch processing, that means it take huge amount of data in input, process it and produce the result. Although batch processing is very efficient for processing high volume of data, but depending on the size of the data being processed and computational power of the system, output can be delayed significantly. Hadoop is not suitable for Real-time data processing.

3. Latency

In Hadoop, MapReduce framework is comparatively slower, since it is designed to support different format, structure and huge volume of data. In MapReduce, Map takes a set of data and converts it into another set of data, where individual element are broken down into key value pair and Reduce takes the output from the map as input and process further and MapReduce requires a lot of time to perform these tasks thereby increasing latency.

4. Ease of use

In Hadoop, MapReduce developers need to hand code for each and every operation which makes it very difficult to work. MapReduce has no interactive mode, but add one such as hive and pig, make working with MapReduce a little easier for adopters.

5. Caching

In Hadoop, MapReduce cannot cache the intermediate data in-memory for a further requirement which diminishes the performance of Hadoop

6. Abstraction

Hadoop does not have any type of abstraction so; MapReduce developers need to hand code for each and every operation which makes it very difficult to work

7. Issue with small files

Hadoop is not suited for small data. HDFS – Hadoop Distributed File System lacks the ability to efficiently support the random reading of small files because of its high capacity design.

What Is RDD?

Resilient Distributed Datasets (RDD) is a fundamental data structure of Spark and main logical data unit. It is an immutable distributed collection of objects. Each dataset in RDD is divided into logical partitions, which may be computed on different nodes of the cluster. RDDs can contain any type of Python, Java, or Scala objects, including user-defined classes.

Distributed means, each RDD is divided into multiple partitions. Each of these partitions can reside in memory or stored on disk of different machines in a cluster. RDDs are immutable (Read Only) data structure. You can't change original RDD, but you can always transform it into different RDD with all changes you want. Formally, an RDD is a read-only, partitioned collection of records. RDDs can be created through deterministic operations on either data on stable storage or other RDDs. RDD is a fault-tolerant collection of elements that can be operated on in parallel.

There are two ways to create RDDs – parallelizing an existing collection in your driver program, or referencing a dataset in an external storage system, such as a shared file system, HDFS, HBase, or any data source offering a Hadoop Input Format.

Spark makes use of the concept of RDD to achieve faster and efficient MapReduce operations.

Features of Spark RDD:

1. **In-memory Computation** Spark RDDs have a provision of in-memory computation. It stores intermediate results in distributed memory (RAM) instead of stable storage (disk).
2. **Lazy Evaluations** All transformations in Apache Spark are lazy, in that they do not compute their results right away. Instead, they just remember the transformations applied to some base data set. Spark computes transformations when an action requires a result for the driver program.
3. **Fault Tolerance** Spark RDDs are fault tolerant as they track data lineage information to rebuild lost data automatically on failure. They rebuild lost data on failure using lineage, each RDD remembers how it was created from other datasets (by transformations like a map, join or groupBy) to recreate itself.
4. **Immutability** Data is safe to share across processes. It can also be created or retrieved anytime which makes caching, sharing & replication easy. Thus, it is a way to reach consistency in computations.
5. **Partitioning** Partitioning is the fundamental unit of parallelism in Spark RDD. Each partition is one logical division of data which is mutable. One can create a partition through some transformations on existing partitions.
6. **Persistence** Users can state which RDDs they will reuse and choose a storage strategy for them (e.g., inmemory storage or on Disk).
7. **Coarse-grained Operations** It applies to all elements in datasets through maps or filter or group by operation.
8. **Location-Stickiness** RDDs are capable of defining placement preference to compute partitions. Placement preference refers to information about the location of RDD. The DAGScheduler places the partitions in such a way that task is close to data as much as possible. Thus, speed up computation.

List down few Spark RDD operations and explain each of them

Transformations:

Any function that returns an RDD is a transformation. Transformations are functions that create a new data set from an existing one by passing each data set element through a function and returning a new RDD representing the results.

All transformations in Spark are lazy. They do not compute their results right away. Instead, they just remember the transformations applied to some base data set (e.g. a file). The transformations are only computed when an action requires a result that needs to be returned to the driver program.

Map:

Map will take each row as input and return an RDD for the row. The map function iterates over every line in RDD and split into new RDD. Using map() transformation we take in any function, and that function is applied to every element of RDD.

For example, in RDD {1, 2, 3, 4, 5} if we apply "rdd.map(x=>x+2)" we will get the result as (3, 4, 5, 6, 7).

Flat Map:

flatMap will take an iterable data as input and returns the RDD as the contents of the iterator. The most simple use of flatMap() is to split each input string into words.

Map and flatMap are similar in the way that they take a line from input RDD and apply a function on that line. The key difference between map() and flatMap() is map() returns only one element, while flatMap() can return a list of elements.

For Example: `val data = spark.read.textFile("spark_test.txt").rdd`

```
val flatmapFile = data.flatMap(lines => lines.split(" ")) flatmapFile.foreach(println)
```

In above code, flatMap() function splits each line when space occurs.

Filter

Spark RDD filter() function returns a new RDD, containing only the elements that meet a predicate. It is a narrow operation because it does not shuffle data from one partition to many partitions.

For Example: `val mapFile = data.flatMap(lines => lines.split(" ")).filter(value => value=="spark")`

```
println(mapFile.count())
```

In above code, flatMap function map line into words and then count the word "Spark" using count() Action after filtering lines containing "Spark" from mapFile.

Distinct()

It returns a new dataset that contains the distinct elements of the source dataset. It is helpful to remove duplicate data.

For example, if RDD has elements (Spark, Spark, Hadoop, Flink), then `rdd.distinct()` will give elements (Spark, Hadoop, Flink).

Reduce By Key

`reduceByKey` takes a pair of key and value pairs and combines all the values for each unique key. So the pairs on the same machine with the same key are combined, before the data is shuffled.

For Example:

```
val words = Array("one","two","two","four","five","six","six","eight","nine","ten")
val data = spark.sparkContext.parallelize(words).map(w => (w,1)).reduceByKey(_+_ )
data.foreach(println)
```

The above code will parallelize the Array of String. It will then map each word with count 1, then `reduceByKey` will merge the count of values having the similar key.

Coalesce

To avoid full shuffling of data we use `coalesce()` function. In `coalesce()` we use existing partition so that less data is shuffled. Using this we can cut the number of the partition. Suppose, we have four nodes and we want only two nodes. Then the data of extra nodes will be kept onto nodes which we kept.

For Example:

```
val rdd1 = spark.sparkContext.parallelize(Array("jan","feb","mar","april","may","jun"),3)
val result = rdd1.coalesce(2)
result.foreach(println)
```

The `coalesce()` will decrease the number of partitions of the source RDD to `numPartitions` define in `coalesce` argument.

Actions:

Actions return final results of RDD computations. Actions trigger execution using lineage graph to load the data into original RDD and carry out all intermediate transformations and return the final results to the Driver program or writes it out to the file system When the action is triggered after the result, new RDD is not formed like transformation. Thus, Actions are Spark RDD operations that give non-RDD values.

Collect

collect is used to return all the elements in the RDD. The application of collect() is unit testing where the entire RDD is expected to fit in memory. As a result, it makes easy to compare the result of RDD with the expected result.

For Example:

```
val data = spark.sparkContext.parallelize(Array(('A',1),('b',2),('c',3)))
val data2 = spark.sparkContext.parallelize(Array(('A',4),('A',6),('b',7),('c',3),('c',8)))
val result = data.join(data2).println(result.collect().mkString(","))
```

The join() transformation in above code will join two RDDs on the basis of same key(alphabet). After that collect() action will return all the elements to the dataset as an Array.

Count

count is used to return the number of elements in the RDD.

For example, RDD has values {1, 2, 2, 3, 4, 5, 5, 6} in this RDD “rdd.count()” will give the result 8

For Example:

```
val data = spark.read.textFile("spark_test.txt").rdd
val mapFile = data.flatMap(lines => lines.split(" ")).filter(value => value=="spark")
println(mapFile.count())
```

In above code flatMap() function maps line into words and count the word “Spark” using count() Action after filtering lines containing “Spark” from mapFile.

Take

The action take(n) returns n number of elements from RDD. It tries to cut the number of partition it accesses, so it represents a biased collection. We cannot presume the order of the elements.

For example, consider RDD {1, 2, 2, 3, 4, 5, 5, 6} in this RDD “take (4)” will give result {2, 2, 3, 4}

For Example:

The take(2) Action will return an array with the first 2 elements of the data set defined in the taking argument.

```
val data = spark.sparkContext.parallelize (Array(('k',5),('s',3),('s',4),('p',7),('p',5),('t',8),('k',6)),3)
val group = data.groupByKey().collect()
val twoRec = result.take(2)
twoRec.foreach(println)
```

Top

If ordering is present in our RDD, then we can extract top elements from our RDD using `top()`. Action `top()` use default ordering of data.

For Example:

```
val data = spark.read.textFile("spark_test.txt").rdd
val mapFile = data.map(line => (line,line.length))
val res = mapFile.top(3)
res.foreach(println)
```

The `map()` operation will map each line with its length. And `top(3)` will return 3 records from `mapFile` with default ordering.

Count By Value

The `countByValue()` returns, many times each element occur in RDD.

For example, RDD has values {1, 2, 2, 3, 4, 5, 5, 6} in this RDD “`rdd.countByValue()`” will give the result {(1,1), (2,2), (3,1), (4,1), (5,2), (6,1)}

For Example:

```
val data = spark.read.textFile("spark_test.txt").rdd
val result= data.map(line => (line,line.length)).countByValue()
result.foreach(println)
```

The `countByValue()` action will return a hashmap of (K, Int) pairs with the count of each key.

Foreach

When we have a situation where we want to apply operation on each element of RDD, but it should not return value to the driver. In this case, `foreach()` function is useful.

For Example:

```
val data = spark.sparkContext.parallelize(Array(('k',5),('s',3),('s',4),('p',7) ,('p',5),('t',8),('k',6)),3)
val group = data.groupByKey().collect()
group.foreach(println)
```

The `foreach()` action run a function (`println`) on each element of the dataset group.